

Haoyue Chang  
Batch code< LISUM22>  
Submission date: July 3<sup>rd</sup>, 2023  
Submitted to Deployment on Flask for Data  
Glacier Internship

```
In [23]: import pandas as pd
import numpy as np
import seaborn as sns
import matplotlib as mpl
import matplotlib.pyplot as plt
# For randomized data splitting
from sklearn.model_selection import train_test_split

# To build linear regression model
import statsmodels.api as sm
# To check model performance
from sklearn.metrics import mean_absolute_error, mean_squared_error
import pickle
```

```
In [24]: df = pd.read_csv('/Users/haoyuechang/Desktop/auto-mpg.csv')
#cData = pd.read_csv("auto-mpg.csv")
```

```
: df = pd.read_csv('/Users/haoyuechang/Desktop/auto-mpg.csv')
#cData = pd.read_csv("auto-mpg.csv")
```

```
: #Data processing
df.info()
#398 entries, 9 columns
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 398 entries, 0 to 397
Data columns (total 9 columns):
#   Column                Non-Null Count  Dtype
---  -
0   mpg                   398 non-null   float64
1   cylinders              398 non-null   int64
2   displacement          398 non-null   float64
3   horsepower             398 non-null   object
4   weight                 398 non-null   int64
5   acceleration           398 non-null   float64
6   model year            398 non-null   int64
7   origin                 398 non-null   int64
8   car name               398 non-null   object
dtypes: float64(3), int64(4), object(2)
memory usage: 28.1+ KB
```

```
: #drop car name since it is not useful to analyze
df1 = df.drop(["car name"], axis=1)
```

```

: hpIsDigit = pd.DataFrame(
    df1.horsepower.str.isdigit()
) # if the string is made of digits store True else False

# print the entries where isdigit = False
df1[hpIsDigit["horsepower"] == False]

```

```

:
      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  origin
32  25.0         4         98.0           ?    2046         19.0         71      1
126  21.0         6        200.0           ?    2875         17.0         74      1
330  40.9         4         85.0           ?    1835         17.3         80      2
336  23.6         4        140.0           ?    2905         14.3         80      1
354  34.5         4        100.0           ?    2320         15.8         81      2
374  23.0         4        151.0           ?    3035         20.5         82      1

```

```

: df1= df1.replace("?", np.nan)
df1[hpIsDigit["horsepower"] == False]

```

```

:
      mpg  cylinders  displacement  horsepower  weight  acceleration  model year  origin
32  25.0         4         98.0         NaN    2046         19.0         71      1
126  21.0         6        200.0         NaN    2875         17.0         74      1
330  40.9         4         85.0         NaN    1835         17.3         80      2
336  23.6         4        140.0         NaN    2905         14.3         80      1
354  34.5         4        100.0         NaN    2320         15.8         81      2
374  23.0         4        151.0         NaN    3035         20.5         82      1

```

```
] : df1.median()
```

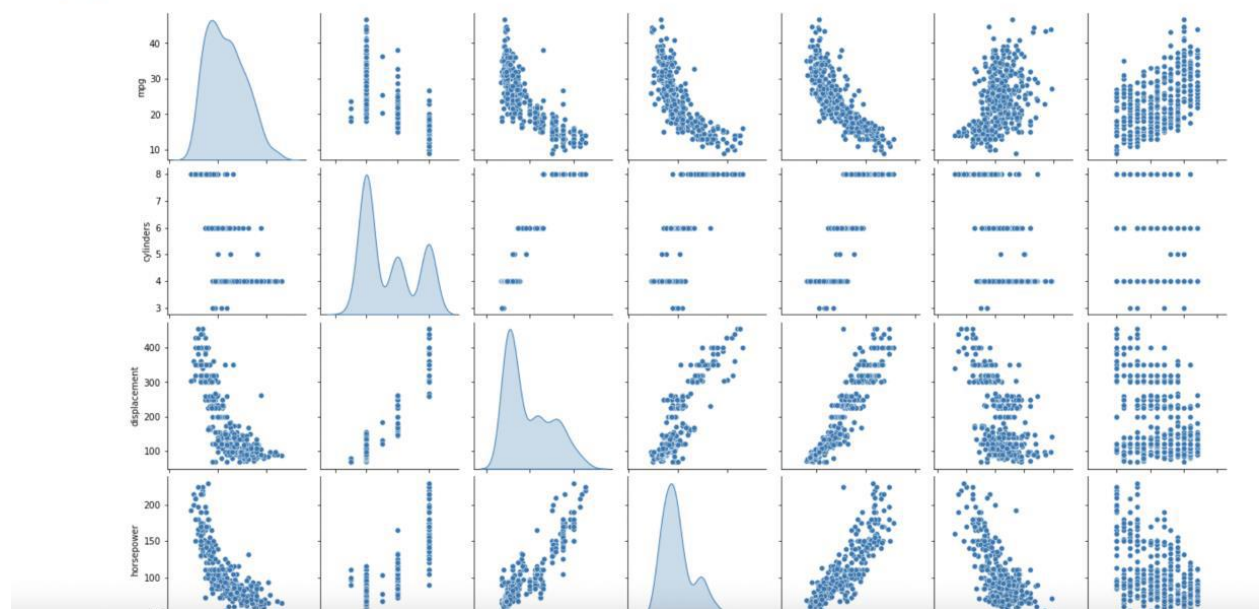
```
]: mpg          23.0
   cylinders     4.0
   displacement  148.5
   horsepower    93.5
   weight       2803.5
   acceleration  15.5
   model year    76.0
   origin        1.0
   dtype: float64
```

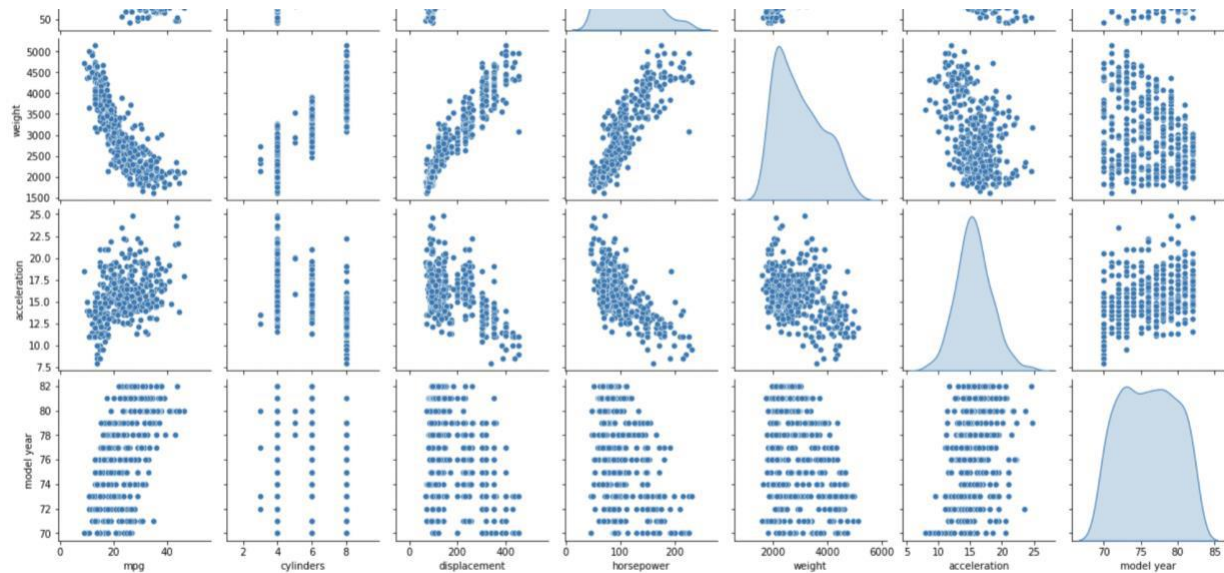
```
] : # Let's replace the missing values with median values of the columns.  
# Note that we do not need to specify the column names below.  
# Every column's missing value is replaced with that column's median respectively  
  
medianFiller = lambda x: x.fillna(x.median())  
df1 = df1.apply(medianFiller, axis=0)
```

```
]: # let's convert the horsepower column from object type to float type
df1["horsepower"] = df1["horsepower"].astype(float)
```

```
In [32]: #Bivariate
df_attr = df1.iloc[:, 0:7]
sns.pairplot(
    df_attr, diag_kind="kde")
```

```
Out[32]: <seaborn.axisgrid.PairGrid at 0x7fb13015ca00>
```





```
In [33]: # drop_first=True will drop one of the three origin columns
#vreat dummy variables
df2 = pd.get_dummies(df1, columns=["origin"], drop_first=True)
df2.head()
```

```
Out[33]:
```

	mpg	cylinders	displacement	horsepower	weight	acceleration	model year	origin_2	origin_3
0	18.0	8	307.0	130.0	3504	12.0	70	0	0
1	15.0	8	350.0	165.0	3693	11.5	70	0	0
2	18.0	8	318.0	150.0	3436	11.0	70	0	0
3	16.0	8	304.0	150.0	3433	12.0	70	0	0
4	17.0	8	302.0	140.0	3449	10.5	70	0	0

```
In [34]: #SPLIT DATA
# independent variables
X = df2.drop(["mpg"], axis=1)
# dependent variable
y = df2[["mpg"]]
```

```
In [35]: X_train, X_test, y_train, y_test = train_test_split(
    X, y, test_size=0.30, random_state=1)
```

```
|: print(X_train.head())
```

	cylinders	displacement	horsepower	weight	acceleration	model year	\
350	4	105.0	63.0	2215	14.9	81	
59	4	97.0	54.0	2254	23.5	72	
120	4	121.0	112.0	2868	15.5	73	
12	8	400.0	150.0	3761	9.5	70	
349	4	91.0	68.0	1985	16.0	81	

	origin_2	origin_3
350	0	0
59	1	0
120	1	0
12	0	0
349	0	1

```
|: print(X_test.head())
```

	cylinders	displacement	horsepower	weight	acceleration	model year	\
174	6	171.0	97.0	2984	14.5	75	
359	4	141.0	80.0	3230	20.4	81	
250	8	318.0	140.0	3735	13.2	78	
274	5	131.0	103.0	2830	15.9	78	
283	6	232.0	90.0	3265	18.2	79	

	origin_2	origin_3
174	0	0
359	1	0
250	0	0
274	1	0
283	0	0

```
|: olsmod = sm.OLS(y_train, X_train)
   olsres = olsmod.fit()
```



```
: print(olsres.summary())
```

```

                        OLS Regression Results
=====
Dep. Variable:          mpg      R-squared (uncentered):          0.980
Model:                  OLS      Adj. R-squared (uncentered):      0.980
Method:                 Least Squares      F-statistic:          1689.
Date:                   Mon, 26 Jun 2023      Prob (F-statistic):      1.04e-225
Time:                   16:45:30      Log-Likelihood:         -741.29
No. Observations:       278      AIC:          1499.
Df Residuals:           270      BIC:          1528.
Df Model:                8
Covariance Type:        nonrobust
=====
                        coef      std err          t      P>|t|      [0.025      0.975]
-----
cylinders              -0.6654      0.427      -1.558      0.120      -1.506      0.175
displacement            0.0284      0.010       2.747      0.006       0.008      0.049
horsepower             -0.0443      0.016      -2.845      0.005      -0.075     -0.014
weight                 -0.0067      0.001      -7.945      0.000      -0.008     -0.005
acceleration           -0.1202      0.110      -1.090      0.277      -0.337      0.097
model_year              0.6229      0.029     21.627      0.000       0.566      0.680
origin_2                2.3585      0.699       3.372      0.001       0.982      3.735
origin_3                2.1483      0.697       3.083      0.002       0.777      3.520
=====
Omnibus:                27.282      Durbin-Watson:          2.241
Prob(Omnibus):           0.000      Jarque-Bera (JB):       42.908
Skew:                    0.609      Prob(JB):               4.82e-10
Kurtosis:                4.490      Cond. No.               1.24e+04
=====
```

#### Notes:

- [1] R<sup>2</sup> is computed without centering (uncentered) since the model does not contain a constant.
- [2] Standard Errors assume that the covariance matrix of the errors is correctly specified.
- [3] The condition number is large, 1.24e+04. This might indicate that there are strong multicollinearity or other numerical problems.

```
# Saving model to disk
pickle.dump(olsres, open('model.pkl', 'wb'))
```

```
# Loading model to compare the results
model = pickle.load(open('model.pkl', 'rb'))
print(model.predict([[8,307.0,130.0,3504,12.0,70,0,0]]))
```

```
[16.27077424]
```

```
!pip3 install flask==2.2.2
```

```
Collecting flask==2.2.2
  Using cached Flask-2.2.2-py3-none-any.whl (101 kB)
Requirement already satisfied: Werkzeug>=2.2.2 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (2.3.6)
Requirement already satisfied: Jinja2>=3.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (3.1.2)
Requirement already satisfied: importlib-metadata>=3.6.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (4.8.1)
Requirement already satisfied: itsdangerous>=2.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (2.1.2)
Requirement already satisfied: click>=8.0 in /opt/anaconda3/lib/python3.9/site-packages (from flask==2.2.2) (8.1.3)
Requirement already satisfied: zipp>=0.5 in /opt/anaconda3/lib/python3.9/site-packages (from importlib-metadata>=3.6.0->flask==2.2.2) (3.6.0)
Requirement already satisfied: MarkupSafe>=2.0 in /opt/anaconda3/lib/python3.9/site-packages (from Jinja2>=3.0->flask==2.2.2) (2.1.2)
Installing collected packages: flask
  Attempting uninstall: flask
    Found existing installation: Flask 2.3.2
    Uninstalling Flask-2.3.2:
      Successfully uninstalled Flask-2.3.2
Successfully installed flask-2.2.2
```



kerrysolar / Heroku-demo

Search Type 🔍 to search

Code

Issues

Pull requests

Actions

Projects

Wiki

Security

Insights

Settings



Heroku-demo

Public

Pin

Unwatch

1

main

1 branch

0 tags

Go to file

Add file

Code



kerrysolar Update Procfile

41c3a25 yesterday 13 commits

templates	Add files via upload	yesterday
LICENSE	Initial commit	yesterday
Procfile	Update Procfile	yesterday
README.md	Initial commit	yesterday
app.py	Update app.py	yesterday
auto-mpg.csv	Add files via upload	yesterday
flask_assignment.ipynb	Add files via upload	yesterday
flask_assignment.py	Add files via upload	yesterday
model.pkl	Add files via upload	yesterday
requirements.txt	Update requirements.txt	yesterday

Enable Automatic Deploys

#### Manual deploy

Deploy the current state of a branch to this app.

#### Deploy a GitHub branch

This will deploy the current state of the branch you specify below. [Learn more](#)

Choose a branch to deploy

main

Deploy Branch

Receive code from GitHub



Build main 41c3a25e



Release phase



Deploy to Heroku



Your app was successfully deployed.

View

## Predict MVG

Cylinders Displacement Horsepower weight Acceleration Model year Origin\_2 Origin\_3 Predict